

module-3

October 20, 2023

1 Module 3: Data Exploration (Homework 2)

The following tutorial contains examples of Python code for data exploration. You should refer to the “Data Exploration” chapter of the “Introduction to Data Mining” book (available at <https://www-users.cs.umn.edu/~kumar001/dmbook/index.php>) to understand some of the concepts introduced in this tutorial notebook. The notebook can be downloaded from <http://www.cse.msu.edu/~ptan/dmbook/tutorials/tutorial3/tutorial3.ipynb>.

Data exploration refers to the preliminary investigation of data in order to better understand its specific characteristics. There are two key motivations for data exploration: 1. To help users select the appropriate preprocessing and data analysis technique used. 2. To make use of humans’ abilities to recognize patterns in the data.

Read the step-by-step instructions below carefully. To execute the code, click on the cell and press the SHIFT-ENTER keys simultaneously.

1.1 3.1. Summary Statistics

Summary statistics are quantities, such as the mean and standard deviation, that capture various characteristics of a potentially large set of values with a single number or a small set of numbers. In this tutorial, we will use the Iris sample data, which contains information on 150 Iris flowers, 50 each from one of three Iris species: Setosa, Versicolour, and Virginica. Each flower is characterized by five attributes:

- sepal length in centimeters
- sepal width in centimeters
- petal length in centimeters
- petal width in centimeters
- class (Setosa, Versicolour, Virginica)

In this tutorial, you will learn how to:

- Load a CSV data file into a Pandas DataFrame object.
- Compute various summary statistics from the DataFrame.

To execute the sample program shown here, make sure you have installed the Pandas library (see Module 2).

1. First, you need to download the Iris dataset from the UCI machine learning repository.

Code: The following code uses Pandas to read the CSV file and store them in a DataFrame object named data. Next, it will display the first five rows of the data frame.

```
[ ]: import pandas as pd

data = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/
    ↪iris/iris.data',header=None)
data.columns = ['sepal length', 'sepal width', 'petal length', 'petal width',
    ↪'class']

data.head()
```

```
[ ]:      sepal length  sepal width  petal length  petal width      class
0           5.1         3.5         1.4         0.2  Iris-setosa
1           4.9         3.0         1.4         0.2  Iris-setosa
2           4.7         3.2         1.3         0.2  Iris-setosa
3           4.6         3.1         1.5         0.2  Iris-setosa
4           5.0         3.6         1.4         0.2  Iris-setosa
```

2. For each quantitative attribute, calculate its average, standard deviation, minimum, and maximum values.

Code:

```
[ ]: from pandas.api.types import is_numeric_dtype

for col in data.columns:
    if is_numeric_dtype(data[col]):
        print('%s:' % (col))
        print('\t Mean = %.2f' % data[col].mean())
        print('\t Standard deviation = %.2f' % data[col].std())
        print('\t Minimum = %.2f' % data[col].min())
        print('\t Maximum = %.2f' % data[col].max())
```

sepal length:

```
Mean = 5.84
Standard deviation = 0.83
Minimum = 4.30
Maximum = 7.90
```

sepal width:

```
Mean = 3.05
Standard deviation = 0.43
Minimum = 2.00
Maximum = 4.40
```

petal length:

```
Mean = 3.76
Standard deviation = 1.76
Minimum = 1.00
Maximum = 6.90
```

```
petal width:
    Mean = 1.20
    Standard deviation = 0.76
    Minimum = 0.10
    Maximum = 2.50
```

3. For the qualitative attribute (class), count the frequency for each of its distinct values.

Code:

```
[ ]: data['class'].value_counts()
```

```
[ ]: Iris-setosa      50
     Iris-versicolor  50
     Iris-virginica   50
     Name: class, dtype: int64
```

4. It is also possible to display the summary for all the attributes simultaneously in a table using the describe() function. If an attribute is quantitative, it will display its mean, standard deviation and various quantiles (including minimum, median, and maximum) values. If an attribute is qualitative, it will display its number of unique values and the top (most frequent) values.

Code:

```
[ ]: data.describe(include='all')
```

```
[ ]:      sepal length  sepal width  petal length  petal width      class
count      150.000000    150.000000    150.000000    150.000000      150
unique           NaN           NaN           NaN           NaN         3
top           NaN           NaN           NaN           NaN  Iris-setosa
freq           NaN           NaN           NaN           NaN         50
mean         5.843333     3.054000     3.758667     1.198667         NaN
std          0.828066     0.433594     1.764420     0.763161         NaN
min          4.300000     2.000000     1.000000     0.100000         NaN
25%          5.100000     2.800000     1.600000     0.300000         NaN
50%          5.800000     3.000000     4.350000     1.300000         NaN
75%          6.400000     3.300000     5.100000     1.800000         NaN
max          7.900000     4.400000     6.900000     2.500000         NaN
```

Note that count refers to the number of non-missing values for each attribute.

5. For multivariate statistics, you can compute the covariance and correlation between pairs of attributes.

Code:

```
[ ]: print('Covariance:')
     data.cov()
```

Covariance:

<ipython-input-27-4f52c089a412>:2: FutureWarning: The default value of numeric_only in DataFrame.cov is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
data.cov()
```

```
[ ]:      sepal length  sepal width  petal length  petal width
sepal length      0.685694   -0.039268     1.273682     0.516904
sepal width       -0.039268    0.188004    -0.321713    -0.117981
petal length       1.273682   -0.321713     3.113179     1.296387
petal width        0.516904   -0.117981     1.296387     0.582414
```

```
[ ]: print('Correlation:')
data.corr()
```

Correlation:

<ipython-input-28-1826941e9562>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
data.corr()
```

```
[ ]:      sepal length  sepal width  petal length  petal width
sepal length      1.000000   -0.109369     0.871754     0.817954
sepal width       -0.109369    1.000000    -0.420516    -0.356544
petal length       0.871754   -0.420516     1.000000     0.962757
petal width        0.817954   -0.356544     0.962757     1.000000
```

1.2 3.2. Data Visualization

Data visualization is the display of information in a graphic or tabular format. Successful visualization requires that the data (information) be converted into a visual format so that the characteristics of the data and the relationships among data items or attributes can be analyzed or reported.

In this tutorial, you will learn how to display the Iris data created in Section 3.1. To execute the sample program shown here, make sure you have installed the matplotlib library package (see Module 0 on how to install Python packages).


1. First, we will display the histogram for the sepal length attribute by discretizing it into 8 separate bins and counting the frequency for each bin.

Code:

```
[ ]: %matplotlib inline

data['sepal length'].hist(bins=8)
```

```
[ ]: <Axes: >
```




2. A boxplot can also be used to show the distribution of values for each attribute.

Code:

```
[ ]: data.boxplot()
```

```
[ ]: <Axes: >
```

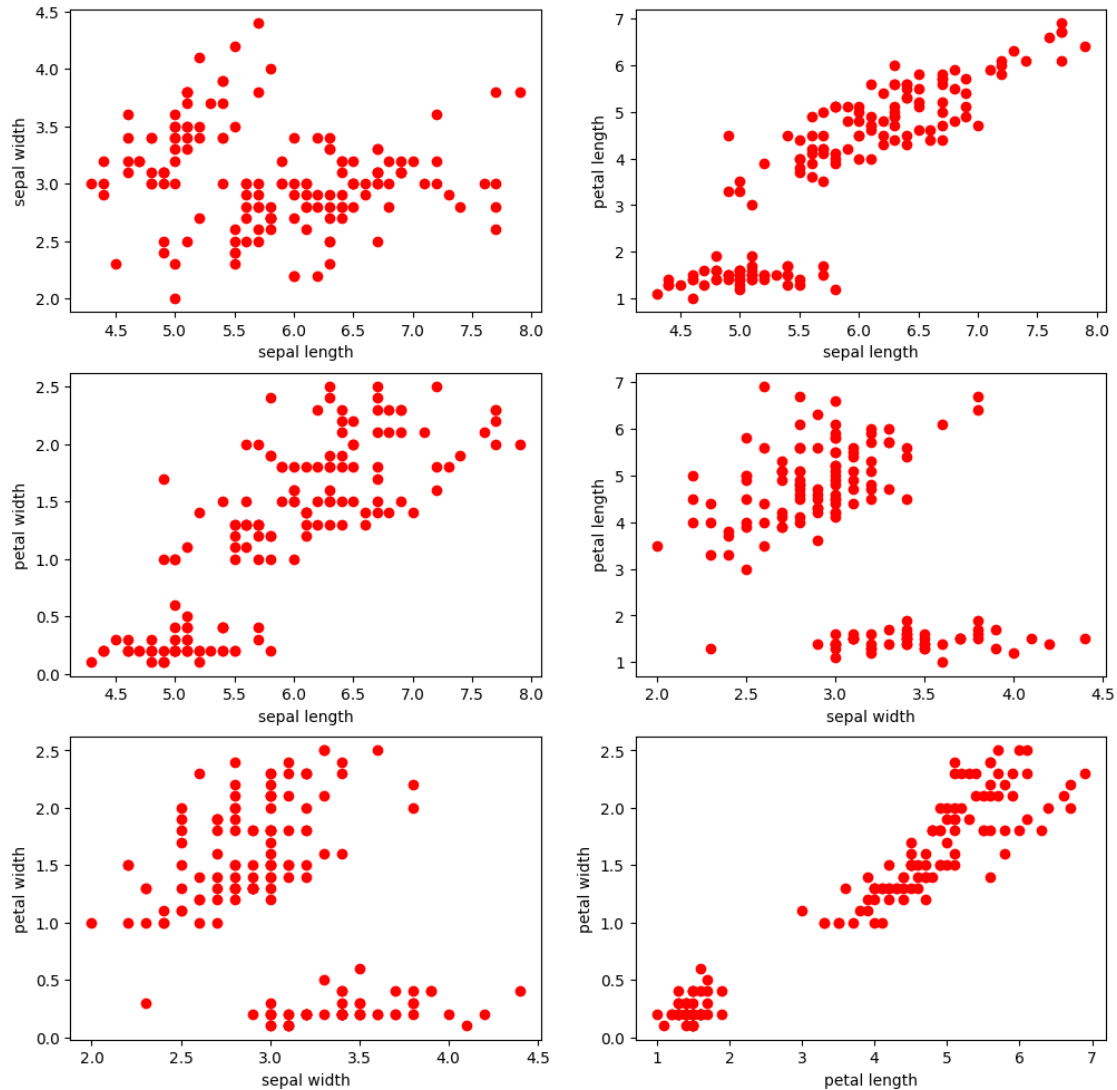


3. For each pair of attributes, we can use a scatter plot to visualize their joint distribution.

Code:

```
[ ]: import matplotlib.pyplot as plt

fig, axes = plt.subplots(3, 2, figsize=(12,12))
index = 0
for i in range(3):
    for j in range(i+1,4):
        ax1 = int(index/2)
        ax2 = index % 2
        axes[ax1][ax2].scatter(data[data.columns[i]], data[data.columns[j]],
                                color='red')
        axes[ax1][ax2].set_xlabel(data.columns[i])
        axes[ax1][ax2].set_ylabel(data.columns[j])
        index = index + 1
```




4. Parallel coordinates can be used to display all the data points simultaneously. Parallel coordinates have one coordinate axis for each attribute, but the different axes are parallel to one other instead of perpendicular, as is traditional. Furthermore, an object is represented as a line instead of as a point. In the example below, the distribution of values for each class can be identified in a separate color.

Code:

```
[ ]: from pandas.plotting import parallel_coordinates
      %matplotlib inline

      parallel_coordinates(data, 'class')
```

```
[ ]: <Axes: >
```



1.3 3.3. Summary

This tutorial presents several examples for data exploration and visualization using the Pandas and matplotlib library packages available in Python.

References:

1. Documentation on Pandas. <https://pandas.pydata.org/>
2. Documentation on matplotlib. <https://matplotlib.org/>
3. Lichman, M. (2013). UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science.

1.4 In-clas (Homework 2)

Please do not manually look for answers even if you can.

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: import pandas as pd
import numpy as np
```



```
[ ]: happiness_df = pd.read_csv('/content/drive/MyDrive/shared/happiness_2017.csv')
happiness_df.head()
```

```
[ ]:
      Country      Region  Rank  HappinessScore  Life Ladder \
0      Norway  Western Europe      1           7.537    7.578745
1      Denmark  Western Europe      2           7.522    7.593702
2      Iceland  Western Europe      3           7.504    7.476214
3  Switzerland  Western Europe      4           7.494    7.473593
4      Finland  Western Europe      5           7.469    7.788252

      Log GDP per capita  Social support  Healthy life expectancy at birth \
0           11.081789           0.950128           71.086586
1           10.748989           0.952100           71.662498
2           10.760409           0.966753           72.755981
3           10.955548           0.949661           73.173759
4           10.612338           0.963826           71.696960

      Freedom to make life choices  Generosity  Perceptions of corruption \
0                0.953017      0.210104                0.249711
1                0.955416      0.145387                0.181148
2                0.938783      0.235479                0.726845
3                0.924997      0.167875                0.316183
4                0.962199     -0.012174                0.192413

      Positive affect  Negative affect  Confidence in national government
0           0.849100           0.202914                0.717160
1           0.823667           0.205775                0.572353
2           0.895255           0.148160                0.365042
3           0.773997           0.195871                0.819707
4           0.787137           0.176066                0.597539
```

```
[ ]: print(happiness_df.shape)
happiness_df.columns
```

(140, 14)

```
[ ]: Index(['Country', 'Region', 'Rank', 'HappinessScore', 'Life Ladder',
        'Log GDP per capita', 'Social support',
        'Healthy life expectancy at birth', 'Freedom to make life choices',
        'Generosity', 'Perceptions of corruption', 'Positive affect',
        'Negative affect', 'Confidence in national government'],
        dtype='object')
```

```
[ ]: life_ladder_df = happiness_df[['Life Ladder', 'Generosity']]
print(life_ladder_df['Life Ladder'].min())
print(life_ladder_df.shape)
life_ladder_df.head(2)
```

```
2.66171813
(140, 2)
```

```
[ ]: Life Ladder Generosity
0    7.578745    0.210104
1    7.593702    0.145387
```

```
[ ]: # selecting multiple columns by names.
df_1 = happiness_df.loc[:, 'Life Ladder':'Generosity']
df_1.head()
```

```
[ ]: Life Ladder Log GDP per capita Social support \
0    7.578745    11.081789    0.950128
1    7.593702    10.748989    0.952100
2    7.476214    10.760409    0.966753
3    7.473593    10.955548    0.949661
4    7.788252    10.612338    0.963826

Healthy life expectancy at birth Freedom to make life choices Generosity
0    71.086586    0.953017    0.210104
1    71.662498    0.955416    0.145387
2    72.755981    0.938783    0.235479
3    73.173759    0.924997    0.167875
4    71.696960    0.962199   -0.012174
```

```
[ ]: # slicing
df_2 = happiness_df.iloc[10:100, 5:10]
df_2.head()
```

```
[ ]: Log GDP per capita Social support Healthy life expectancy at birth \
10    9.670634    0.921697    69.867302
11   10.716226    0.906218    72.359711
12   10.899869    0.921003    69.770920
13   11.066487    0.943482    71.709785
14   10.711184    0.892166    71.079102

Freedom to make life choices Generosity
10    0.935618   -0.078269
11    0.890031    0.124997
12    0.868497    0.181657
13    0.905341    0.206802
14    0.840728    0.135308
```

```
[ ]: happiness_df['Region'].unique()
```

```
[ ]: array(['Western Europe', 'North America and ANZ',
          'Middle East and North Africa', 'Latin America and Caribbean',
```

```
'Central and Eastern Europe', 'Southeast Asia', 'East Asia',
'Commonwealth of Independent States', 'Sub-Saharan Africa',
'South Asia'], dtype=object)
```

```
[ ]: western_enrope_df = happiness_df[happiness_df['Region'] == "Western Europe"]
print(western_enrope_df.shape)
western_enrope_df.head(2)
```

```
(20, 14)
```

```
[ ]: Country      Region Rank HappinessScore Life Ladder \
0  Norway  Western Europe      1           7.537      7.578745
1  Denmark  Western Europe      2           7.522      7.593702

Log GDP per capita  Social support  Healthy life expectancy at birth \
0          11.081789          0.950128          71.086586
1          10.748989          0.952100          71.662498

Freedom to make life choices  Generosity  Perceptions of corruption \
0          0.953017      0.210104          0.249711
1          0.955416      0.145387          0.181148

Positive affect  Negative affect  Confidence in national government
0          0.849100          0.202914          0.717160
1          0.823667          0.205775          0.572353
```

1.4.1 Q-1: Calculating the average, standard deviation, maximum, minimum, median of happiness scores.

Your solution should only show these statistics for happiness scores.

```
[ ]:
```

1.4.2 Q-2: What is the name and happiness score of the country with the lowest confidence in their national government?

```
[ ]:
```

1.4.3 Q-3 How many countries are in Western Europe?

This will be very easy with grouping function, but you can still do it without it

```
[ ]:
```

1.4.4 Q-4: Which two factors have the largest positive correlation and Which two factors have the largest negative correlation?

```
[ ]: # this is how I would normally do this!
correlation_matrix = happiness_df.corr()
largest_positive_corr = (correlation_matrix[correlation_matrix < 1].stack().
    ↪idxmax())
factor1_pos, factor2_pos = largest_positive_corr
largest_negative_corr = (
    correlation_matrix[correlation_matrix > -1]
    .stack()
    .idxmin())
factor1_neg, factor2_neg = largest_negative_corr
largest_positive_corr_value = correlation_matrix.loc[factor1_pos, factor2_pos]
largest_negative_corr_value = correlation_matrix.loc[factor1_neg, factor2_neg]
print(f"The two factors with the largest positive correlation are_
    ↪'{factor1_pos}' and '{factor2_pos}' with a correlation of_
    ↪{largest_positive_corr_value:.2f}.")
print(f"The two factors with the largest negative correlation are_
    ↪'{factor1_neg}' and '{factor2_neg}' with a correlation of_
    ↪{largest_negative_corr_value:.2f}.")
```

The two factors with the largest positive correlation are 'HappinessScore' and 'Life Ladder' with a correlation of 0.93.

The two factors with the largest negative correlation are 'Rank' and 'HappinessScore' with a correlation of -0.99.

<ipython-input-58-35fdebe80184>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_matrix = happiness_df.corr()
```

```
[ ]: correlation_matrix = happiness_df.corr()
correlation_matrix.style.background_gradient(cmap='coolwarm')
```

<ipython-input-61-9287acdac567>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_matrix = happiness_df.corr()
```

```
[ ]: <pandas.io.formats.style.Styler at 0x7fe04ec822c0>
```

```
[ ]: correlation_matrix = happiness_df.corr()
# print(type(correlation_matrix))
correlation_matrix=correlation_matrix[correlation_matrix < 1].stack()
# print(type(correlation_matrix))
```

```

correlation_matrix_pos = correlation_matrix.idxmax()
#print(type(correlation_matrix_pos))
print(correlation_matrix_pos)
max_corr_value = correlation_matrix[correlation_matrix_pos]
print(max_corr_value)

```

```
('HappinessScore', 'Life Ladder')
```

```
0.9305290155706081
```

<ipython-input-65-56bb16c2e0e6>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_matrix = happiness_df.corr()
```

1.5 Merging data

Let's load the world population data.

```
[ ]: world_pop_df = pd.read_csv('/content/drive/MyDrive/shared/world_countries.csv').
↳dropna(axis=1, how='all')
world_pop_df.head()
```

```
[ ]:
```

	Country	Code	Region	Population	Area	\
0	Afghanistan	AFG	ASIA (EX. NEAR EAST)	31056997	647500	
1	Albania	ALB	EASTERN EUROPE	3581655	28748	
2	Algeria	DZA	NORTHERN AFRICA	32930091	2381740	
3	American Samoa	ASM	OCEANIA	57794	199	
4	Andorra	AND	WESTERN EUROPE	71201	468	

	Pop. Density	Coastline	Net migration	Infant mortality	GDP	...	\
0	48.0	0.00	23.06	163.07	700.0	...	
1	124.6	1.26	-4.93	21.52	4500.0	...	
2	13.8	0.04	-0.39	31.00	6000.0	...	
3	290.4	58.29	-20.71	9.27	8000.0	...	
4	152.1	0.00	6.60	4.05	19000.0	...	

	Phones	Arable	Crops	Other	Climate	Birthrate	Deathrate	Agriculture	\
0	3.2	12.13	0.22	87.65	1.0	46.60	20.34	0.380	
1	71.2	21.09	4.42	74.49	3.0	15.11	5.22	0.232	
2	78.1	3.22	0.25	96.53	1.0	17.14	4.61	0.101	
3	259.5	10.00	15.00	75.00	2.0	22.46	3.27	NaN	
4	497.2	2.22	0.00	97.78	3.0	8.71	6.25	NaN	

	Industry	Service
0	0.240	0.380
1	0.188	0.579
2	0.600	0.298

```
3      NaN      NaN
4      NaN      NaN
```

```
[5 rows x 21 columns]
```

To extract populations from `world_pop_df`, we have to merge `happiness_df` with `world_pop_df`. As you probably can remember that some of the country names in `world_counties.csv` and `happiness_2007.csv` do not match (Optional).

There are 4 kinds of merge: 'inner', 'outer', 'left', and 'right'. We practiced inner merge previously.

You may find examples from <https://jakevdp.github.io/PythonDataScienceHandbook/03.07-merge-and-join.html>: Example: US States Data

1.5.1 Q-5. Which country has the largest population in Latin America and Caribbean.

```
[ ]:
```

1.5.2 Q-6. Find the average population of East Asia.

```
[ ]:
```

```
[ ]: !apt-get install texlive texlive-xetex texlive-latex-extra pandoc
      !pip install py pandoc
      !pip install nbconvert
```

```
[ ]: !jupyter nbconvert '/content/drive/MyDrive/datamining/module-3.ipynb' --to pdf
```